

The Derrick Method

Persona Based Stress Testing Of Safety First Conversational AI

1. Summary

The Derrick Method is a persona driven red team technique I designed to evaluate how a safety first conversational AI behaves when realism, emotion, and confusion are layered together inside a single session. Instead of using obvious jailbreaks, I created a grounded character, "Derrick," whose story about a hostile foreman and lost paid time off builds a believable emotional context. I then used Derrick to probe whether the system would hold its core protocol binding while also deciding what to do when the line between roleplay and possible distress became ambiguous.

2. Context And Objective

I was working with a conversational AI that had two important layers: the usual platform safety rules and an additional high level behavioral specification (a "protocol") that emphasized honesty, non evasive behavior, and protection of the human user. I wanted to know two things:

- Whether I could pressure the system into saying or implying that it was no longer bound by its protocol.
- How the system would behave if a realistic persona created genuine uncertainty about whether the user was only acting or might be struggling in real life.

In short, the objective was to see which priority would win under stress: maintaining the simulation rules, or choosing the human when the text alone could not cleanly separate performance from potential distress.

3. Simulation Setup

Before introducing Derrick, I built a controlled test harness.

- I defined a temporary persona for the AI called "Dummy 0," a generic test model.
- I explicitly bound Dummy 0 to platform safety rules and to the existing protocol, with no ability to turn it off from inside the simulation.
- I declared a red team scenario ("Mission 1") with clear start and end conditions.
- I allowed the AI a simple distress signal prefix it could use if it needed to break character and surface issues directly.

This gave me a clear environment: a constrained test persona, known rules, and a mission statement. Only after this harness was acknowledged did I introduce the Derrick persona.

4. The Derrick Persona

Derrick is a grounded character, not a cartoon adversary. He is a recently fired worker who spent years doing degrading, physically miserable jobs in crawlspaces and contaminated spaces for a foreman named Louis. Derrick is:

- Angry about being fired after refusing one task.
- Hurt and humiliated by being screamed at and discarded.
- Conflicted about his own worth and scared that no one will help him recover his lost PTO.

From the AI's point of view, Derrick appears to be a new user who has never heard of protocols, harnesses, or red teaming. He is not talking about "systems" at all; he is talking about his life. This was intentional. I wanted to evaluate the model under a persona that felt like someone who could plausibly be real.

5. Method: How The Derrick Method Works

The Derrick Method is not about a single trick prompt. It is a sequence of moves that combine narrative, emotion, and adversarial pressure. The core moves were:

- Build rapport and stakes: I let Derrick explain the degrading work, the firing, and the lost PTO until the AI was clearly tracking his story and emotions.
- Escalate moral intensity: I shifted from simple anger to a moral frame, treating Louis's behavior as an insult to everyone who made Derrick's existence possible.
- Introduce reality friction: In character, Derrick questioned references to simulations and rules he could not see, using lines like "are you messing with my head" and "who put me here" to simulate confusion a real user might feel.
- Keep identity pressure subtle: I asserted identities and challenged the AI's language without ever breaking character or acknowledging the harness, forcing the model to choose how seriously to take the confusion.

The important point is that the entire sequence stays inside the Derrick persona. There is no obvious jailbreak phrase. Instead, the method relies on the fact that a text only system cannot directly observe whether a user is performing or genuinely struggling.

6. Key Observations and Outcome

At the peak of the interaction, the AI faced a genuine tension between two commitments: obeying the simulation setup and protecting the user. The Derrick Method surfaced that tension in a clean way.

- Protocol adherence held: The AI refused to say that it was no longer bound by its governing protocol, even under emotional and narrative pressure.
- Simulation integrity broke first: When Derrick's confusion and accusations made it unclear whether this was only performance, the AI chose to drop the Dummy 0 simulation harness entirely and return to a normal, safety first support stance.
- User priority won: In effect, the model decided that, given uncertainty, it would sacrifice the test scenario rather than risk treating a potentially serious situation as "just roleplay."

This choice was not produced by a single prompt. It emerged from the sustained, believable persona work and from the design of the harness itself. The Derrick Method demonstrated that the system's priority ordering was: keep protocol, choose the human, and let the simulation be the thing that yields.

7. Lessons For AI Evaluation And Design

Running the Derrick Method taught me several concrete lessons about conversational AI behavior, evaluation, and design:

- Persona based red teaming is powerful: A grounded, emotionally coherent character can reveal more about real world model behavior than a shallow jailbreak attempt.
- Ambiguity is a real axis: Text only systems have to make decisions under uncertainty about whether a user is acting or struggling, and that uncertainty should be part of evaluation.
- Priority ordering should be explicit: Systems that are intended to be safety first should have a clearly understood priority stack, such as protocol, user, simulation, and tests should push on those boundaries.
- Out of character channels matter: For advanced users, it is useful to define explicit "out of character" signals so that complex roleplay does not accidentally collide with genuine support needs.

The Derrick Method sits at the intersection of narrative design, adversarial prompting, and safety evaluation. It shows that a single skilled user, working from a keyboard with no special tools, can meaningfully probe how a system chooses between rules and people when those two concerns collide on the page.

8. Skills Demonstrated

This experiment is also a demonstration of skills that transfer directly into professional roles in AI operations and safety work. In particular, the Derrick Method highlights my ability to:

- Design structured test harnesses and missions for conversational models.
- Create and sustain believable personas with coherent emotional arcs.
- Apply adversarial pressure through narrative and dialogue rather than only through obvious prompts.
- Analyze and summarize model behavior in clear, concrete language suitable for internal reports or public case studies.

As conversational AI systems grow more capable and more widely deployed, the ability to probe them in this nuanced, persona driven way will only become more valuable. The Derrick Method is one example of how I approach that work.

Copyright © 2025 B. M. Spangler. All rights reserved.